



AwesomeNeuron

Newsletter: <https://awesomeneuron.substack.com/>

Linkedin:: <https://www.linkedin.com/in/analyticalrohit>

Code: https://github.com/analyticalrohit/pytorch_fundamentals

PyTorch Fundamentals: Your First Steps into Hands-on Deep Learning

This notebook provides an introduction to PyTorch, covering tensor initialization, operations, indexing, and reshaping. Follow along to learn the basics with clear examples and detailed explanations.

Table of Contents

- What are Tensors?
- Tensor Initialization
- Common Tensor Initialization Methods
- Tensor Type Conversion
- Converting Between NumPy Arrays and Tensors
- Tensor Mathematics and Comparison Operations
- Matrix Multiplication and Batch Operations
- Broadcasting and Other Useful Operations
- Tensor Indexing
- Tensor Reshaping

```
In [22]: import torch
import numpy as np
# Ignore warnings
import warnings
warnings.filterwarnings('ignore')


# Print versions
print("torch version:", torch.__version__)
print("numpy version:", np.__version__)

torch version: 2.4.0+cu118
numpy version: 1.26.3
```



What are Tensors?

Tensor holds a multi-dimensional array of elements of a single data type which is very similar with numpy's ndarray. When the dimension is zero, it can be called a scalar. When the dimension is 2, it can be called a matrix. When the dimension is greater than 2, it is usually called a tensor. If you are very familiar with numpy, then understanding Tensor will be quite easy.



Tensor Initialization

This code creates a 2x3 PyTorch tensor with float32 data type, assigns it to a specified device (CPU or GPU), and enables gradient tracking for backpropagation.

```
In [23]: # Check for CUDA availability and set the device
device = "cuda" if torch.cuda.is_available() else "cpu"

# Initialize a 2x3 tensor with requires_grad enabled
my_tensor = torch.tensor([[1, 2, 3], [4, 5, 6]], dtype=torch.float32, device=device)

print(my_tensor)
print("Data type:", my_tensor.dtype)
print("Device:", my_tensor.device)
print("Shape:", my_tensor.shape)
print("Requires Gradient:", my_tensor.requires_grad)
```

```
tensor([[1., 2., 3.],
        [4., 5., 6.]], device='cuda:0', requires_grad=True)
Data type: torch.float32
Device: cuda:0
Shape: torch.Size([2, 3])
Requires Gradient: True
```

Other Common Tensor Initialization Methods

- **Empty Tensor**: Creates an uninitialized 3x3 tensor (random values).
- **Zeros Tensor**: Creates a 3x3 tensor filled with zeros.

- **Random Tensor:** Generates a 3×3 tensor with random values between 0 and 1.
- **Ones Tensor:** Creates a 3×3 tensor filled with ones. **Identity Matrix:** Generates a 4×4 identity matrix (diagonal of ones).
- **Arange Tensor:** Creates a 1D tensor with values from 0 to 4 (step of 1).
- **Linspace Tensor:** Generates 5 evenly spaced values between 0.1 and 1.
- **Normal Distributed Tensor:** Fills a tensor with values from a normal (Gaussian) distribution with mean 0 and std 1.
- **Uniform Distributed Tensor:** Fills a tensor with values from a uniform distribution between 0 and 1.
- **Diagonal Tensor:** Creates a 4×4 diagonal tensor with ones along the diagonal and zeros elsewhere.

```
In [24]: # Create an empty tensor of size 3x3
x = torch.empty(3, 3)
print("Empty Tensor:\n", x)

# Create a tensor filled with zeros
x = torch.zeros(3, 3)
print("Zeros Tensor:\n", x)

# Create a tensor with random values
x = torch.rand(3, 3)
print("Random Tensor:\n", x)

# Create a tensor filled with ones
x = torch.ones(3, 3)
print("Ones Tensor:\n", x)

# Create an identity matrix
x = torch.eye(4, 4)
print("Identity Matrix:\n", x)

# Create a tensor using arange
x = torch.arange(5)
print("Arange Tensor:\n", x)

# Create a tensor using linspace
x = torch.linspace(0.1, 1, 5)
print("Linspace Tensor:\n", x)

# Create a tensor with values drawn from a normal distribution
x = torch.empty(1, 5).normal_(mean=0, std=1)
print("Normal Distributed Tensor:\n", x)

# Create a tensor with values drawn from a uniform distribution
x = torch.empty(1, 5).uniform_(0, 1)
print("Uniform Distributed Tensor:\n", x)
```



```

Empty Tensor:
tensor([[1., 1., 1.],
        [1., 1., 1.],
        [1., 1., 1.]])
Zeros Tensor:
tensor([[0., 0., 0.],
        [0., 0., 0.],
        [0., 0., 0.]])
Random Tensor:
tensor([[0.9643, 0.1459, 0.7576],
        [0.3014, 0.5924, 0.1967],
        [0.5843, 0.1187, 0.0444]])
Ones Tensor:
tensor([[1., 1., 1.],
        [1., 1., 1.],
        [1., 1., 1.]])
Identity Matrix:
tensor([[1., 0., 0., 0.],
        [0., 1., 0., 0.],
        [0., 0., 1., 0.],
        [0., 0., 0., 1.]])
Arange Tensor:
tensor([0, 1, 2, 3, 4])
Linspace Tensor:
tensor([0.1000, 0.3250, 0.5500, 0.7750, 1.0000])
Normal Distributed Tensor:
tensor([[-0.7517, -0.4783, -1.6978,  1.0087, -0.1160]])
Uniform Distributed Tensor:
tensor([[0.0408, 0.3164, 0.1182, 0.5151, 0.5141]])

```

Tensor Type Conversion

Creates a tensor with values [0, 1, 2, 3] and demonstrates type conversion to boolean, int16, int64, float16, float32, and float64.

```

In [25]: # Create a tensor and convert its type
tensor = torch.arange(4)
print("Boolean Tensor:", tensor.bool())    # Convert to boolean
print("Short Tensor (int16):", tensor.short()) # Convert to int16
print("Long Tensor (int64):", tensor.long())  # Convert to int64
print("Half Tensor (float16):", tensor.half()) # Convert to float16
print("Float Tensor (float32):", tensor.float()) # Convert to float32
print("Double Tensor (float64):", tensor.double()) # Convert to float64

```

```

Boolean Tensor: tensor([False,  True,  True,  True])
Short Tensor (int16): tensor([0, 1, 2, 3], dtype=torch.int16)
Long Tensor (int64): tensor([0, 1, 2, 3])
Half Tensor (float16): tensor([0., 1., 2., 3.], dtype=torch.float16)
Float Tensor (float32): tensor([0., 1., 2., 3.])
Double Tensor (float64): tensor([0., 1., 2., 3.], dtype=torch.float64)

```

Converting Between NumPy Arrays and Tensors

PyTorch makes it easy to switch between NumPy arrays and tensors, allowing seamless integration with existing computing workflows.

```
In [26]: # Create a NumPy array of zeros
np_array = np.zeros((5, 5))
print("NumPy Array:\n", np_array)

# Convert NumPy array to PyTorch tensor
tensor = torch.from_numpy(np_array)
print("Tensor from NumPy Array:\n", tensor)

# Convert tensor back to NumPy array
numpy_back = tensor.numpy()
print("Converted Back to NumPy Array:\n", numpy_back)
```

NumPy Array:

```
[[0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0.]]
```

Tensor from NumPy Array:

```
tensor([[0., 0., 0., 0., 0.],
        [0., 0., 0., 0., 0.],
        [0., 0., 0., 0., 0.],
        [0., 0., 0., 0., 0.],
        [0., 0., 0., 0., 0.]], dtype=torch.float64)
```

Converted Back to NumPy Array:

```
[[0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0.]]
```

Tensor Mathematics and Comparison Operations

This section explores essential math operations with PyTorch tensors.

- **Addition & Subtraction:** Adds and subtracts two tensors element-wise.
- **Division:** Uses true division for precise results.
- **Inplace Operations:** Modifies a tensor directly without creating a new one.
- **Exponentiation:** Raises each element to a power using `pow()` or `**`.
- **Comparisons:** Checks conditions like `x > 0` or `x < 0`, returning boolean results.
- **Dot Product:** Computes the sum of element-wise multiplications between two tensors.



```
In [27]: # Define two tensors for operations
x = torch.tensor([1, 2, 3]) y =
torch.tensor([9, 8, 7])

# Addition
z = x + y
print("Addition Results:", z)

# Addition using .add
z1 = torch.empty(3)
torch.add(x, y, out=z1)
z2 = torch.add(x, y)
print("Addition Results:", z, z1, z2)

# Subtraction
z = x - y
print("Subtraction Result:", z)

# Division (true division)
z = torch.true_divide(x, y)
print("Division Result:", z)

# Inplace operations
t = torch.ones(3)
print("Before inplace addition:", t)
t.add_(x)
print("After inplace addition:", t)
t += x # Another inplace addition (note: t = t + x creates a new tensor)
print("After second inplace addition:", t)

# Exponentiation
z = x.pow(2)
print("Exponentiation (pow):", z)
z = x**2
print("Exponentiation (**):", z)

# Comparisons
z = x > 0
print("x > 0:", z)
z = x < 0
print("x < 0:", z)

# Dot product
z = torch.dot(x, y)
print("Dot Product:", z)
```



```

Addition Results: tensor([10, 10, 10])
Addition Results: tensor([10, 10, 10]) tensor([10., 10., 10.]) tensor([10, 1
0, 10])
Subtraction Result: tensor([-8, -6, -4])
Division Result: tensor([0.1111, 0.2500, 0.4286])
Before inplace addition: tensor([1., 1., 1.])
After inplace addition: tensor([2., 3., 4.])
After second inplace addition: tensor([3., 5., 7.])
Exponentiation (pow): tensor([1, 4, 9])
Exponentiation (**): tensor([1, 4, 9])
x > 0: tensor([True, True, True])
x < 0: tensor([False, False, False])
Dot Product: tensor(46)

```

Matrix Multiplication and Batch Operations

Matrix operations are at the heart of deep learning. Let's find out different ways to perform multiplication.

- **Matrix Multiplication:** Uses `@` or `torch.mm()` to perform standard matrix multiplication.
- **Matrix Exponentiation:** Raises a square matrix to a power using `matrix_power(n)`.
- **Element-wise Multiplication:** Uses `torch.mul()` or `*` for element-wise multiplication.
- **Batch Matrix Multiplication:** Uses `torch.bmm()` to multiply batches of matrices efficiently.

```

In [28]: # Matrix multiplication using @ operator and torch.mm
x2 = torch.tensor([[1, 2, 3]]) y2 = torch.tensor([[9,
8, 7]])

z = x2 @ torch.t(y2)
print("Matrix Multiplication (@ operator):\n", z)
z = torch.mm(x2, torch.t(y2))
print("Matrix Multiplication (torch.mm):\n", z)
z = x2.mm(torch.t(y2))
print("Matrix Multiplication (mm):\n", z)

# Matrix exponentiation: multiplying a matrix with itself 3 times
matrix_exp = torch.rand(5, 5)
print("Matrix multiplied 3 times:\n", matrix_exp @ matrix_exp @ matrix_exp)
print("Matrix power 3:\n", matrix_exp.matrix_power(3))

# Element-wise multiplication
z = torch.mul(x, y)
print("Element-wise Multiplication:", z)
z = x * y
print("Element-wise Multiplication (alternative):", z)

# Batch matrix multiplication

```

```
batch = 32 n, m, p = 10, 20, 30
tensor1 = torch.rand((batch, n, m))
tensor2 = torch.rand((batch, m, p))
out_bmm = torch.bmm(tensor1, tensor2)

# Result shape: (batch, n, p)
print("Batch Matrix Multiplication (first batch):\n", out_bmm[0])
print("Shape of batched multiplication result:", (tensor1 @ tensor2).shape)
```



```

Matrix Multiplication (@ operator):
tensor([[46]])
Matrix Multiplication (torch.mm):
tensor([[46]])
Matrix Multiplication (mm):
tensor([[46]])
Matrix multiplied 3 times:
tensor([[6.8649, 7.1375, 6.5783, 7.6163, 6.8593],
        [6.4644, 6.5831, 6.2444, 7.0148, 6.4003],
        [5.8586, 6.0872, 5.5466, 6.2068, 5.5633],
        [6.6579, 6.8326, 6.3734, 7.4843, 6.8450],
        [5.6023, 5.7104, 5.4097, 5.7048, 5.1312]])
Matrix power 3:
tensor([[6.8649, 7.1375, 6.5783, 7.6163, 6.8593],
        [6.4644, 6.5831, 6.2444, 7.0148, 6.4003],
        [5.8586, 6.0872, 5.5466, 6.2068, 5.5633],
        [6.6579, 6.8326, 6.3734, 7.4843, 6.8450],
        [5.6023, 5.7104, 5.4097, 5.7048, 5.1312]])
Element-wise Multiplication: tensor([ 9, 16, 21]) Element-wise Multiplication
(alternative): tensor([ 9, 16, 21]) Batch Matrix Multiplication (first
batch): tensor([[3.5440, 6.5026, 6.5445, 4.5252, 5.2045, 5.7321, 6.1181,
6.6445, 6. 0093,

858,      4.9353, 3.9677, 4.6019, 6.3319, 5.7043, 5.1247, 4.3091, 4.7326, 4.8
894,      5.1322, 5.6088, 5.9398, 6.9429, 5.9886, 5.2573, 2.7003, 4.8218, 5.9
      4.6512, 5.1542, 3.7731],
781, [4.2282, 5.0751, 4.5581, 3.9363, 5.2671, 5.2506, 4.4924, 5.5095, 5.1
779,      3.7145, 4.7497, 3.8810, 4.2245, 4.6506, 4.5174, 2.8565, 3.9761, 3.8
386,      4.3638, 4.0318, 5.5015, 3.9265, 5.3370, 4.7982, 2.2797, 4.3291, 3.7
      4.0959, 3.6940, 3.2689],
909, [4.5827, 6.1946, 4.9569, 4.6735, 5.6490, 5.8873, 5.7623, 6.4767, 6.3
429,      4.4470, 5.0609, 4.6619, 5.1481, 5.3259, 5.4949, 3.7388, 5.1354, 4.2
472,      4.5741, 4.5234, 5.5854, 4.6225, 5.9995, 5.5425, 2.8380, 4.2748, 4.3
      4.4736, 4.1825, 3.4843],
801, [3.3100, 5.1634, 5.0415, 3.7873, 4.4239, 5.4084, 5.0781, 5.1528, 5.1
439,      3.5463, 3.6643, 3.9544, 4.7713, 4.8103, 3.8891, 2.8634, 4.0375, 3.4
179,      4.1778, 4.0538, 4.9986, 4.8988, 5.2553, 4.8321, 1.6977, 4.3844, 4.1
      4.2077, 4.1596, 3.0347],
557, [4.1352, 5.5097, 5.4081, 4.1339, 5.2045, 5.8074, 5.9689, 5.7145, 5.1
134,      4.5157, 4.7473, 5.3542, 5.0027, 4.4979, 5.2767, 4.4354, 3.5934, 4.7
712,      5.1901, 5.5433, 6.2373, 4.5942, 5.5256, 5.3336, 3.2117, 4.1416, 4.2

```



```

        4.1048, 4.0465, 3.6511],
    [3.7322, 5.3081, 5.4181, 4.6393, 5.1372, 5.5894, 5.5543, 6.0296, 4.5
670,
        4.0552, 4.2119, 4.5391, 5.2059, 4.8795, 5.0137, 3.7978, 3.7304, 4.0
530,
        4.3719, 5.0318, 5.7526, 4.8653, 4.8926, 5.5795, 2.6404, 4.2889, 4.5
551,
        4.2022, 3.9910, 3.8555],
    [3.2315, 4.6632, 4.9304, 3.5092, 4.4533, 4.8799, 4.6290, 5.3459, 4.9
640,
        3.7996, 4.0911, 4.0296, 4.6814, 4.0135, 4.5876, 3.6069, 3.8043, 4.2
701,
        4.1166, 4.8829, 5.5559, 4.8536, 4.5921, 3.6704, 2.4689, 3.4953, 3.6
832,
        3.9577, 3.5755, 3.0293],
    [4.7066, 5.7029, 6.1439, 4.4224, 5.7423, 6.4657, 5.8944, 5.8826, 6.2
143,
        4.1787, 4.8364, 4.9471, 5.2735, 4.8664, 5.5245, 3.1057, 4.8749, 4.4
022,
        5.6390, 5.2407, 6.2896, 5.6143, 5.5157, 5.6265, 2.1874, 4.4567, 3.9
869,
        5.5333, 4.4775, 3.9314],
    [4.3227, 5.8580, 5.9357, 4.1483, 5.2436, 6.3560, 5.7144, 5.9563, 5.8
249,
        4.5020, 4.4281, 4.4165, 5.4092, 4.8056, 4.9667, 4.0844, 4.2222, 3.8
133,
        4.5904, 5.0534, 6.7042, 5.0727, 5.5185, 5.6069, 3.0032, 4.5262, 4.0
067,
        4.7430, 4.2932, 3.7362],
    [3.8294, 5.4057, 5.0970, 3.9448, 5.0699, 5.1357, 4.9649, 5.2615, 4.4
683,
        3.5833, 4.0856, 3.3753, 4.0069, 4.1184, 4.2750, 3.4178, 3.7880, 3.3
670,
        4.7347, 4.6245, 5.3556, 3.9735, 5.1346, 5.1036, 2.6921, 4.2103, 3.9
386,
        3.6596, 4.1051, 2.2051]])
Shape of batched multiplication result: torch.Size([32, 10, 30])

```

Broadcasting and Other Useful Operations

Broadcasting allows arithmetic operations on tensors of different shapes. This section also demonstrates additional useful functions.

- **Broadcasting:** Automatically expands smaller tensors to match larger ones in operations.
- **Summation:** `torch.sum(x, dim=0)` computes sum along a specific dimension.
- **Min/Max Values:** `torch.max()` and `torch.min()` return the highest and lowest values along a dimension.
- **Absolute Values:** `torch.abs(x)` gets the element-wise absolute values.



- **Argmax/Argmin:** `torch.argmax()` and `torch.argmin()` return the index of max/min values.
- **Mean Calculation:** `torch.mean(x.float(), dim=0)` computes the mean (ensuring float dtype).
- **Element-wise Comparison:** `torch.eq(x, y)` checks equality between two tensors.
- **Sorting:** `torch.sort(y, dim=0)` sorts tensor elements and returns indices.
- **Clamping:** `torch.clamp(x, min=0)` restricts values within a range.
- **Boolean Operations:** `torch.any(x_bool)` checks if any value is True, `torch.all(x_bool)` checks if all are True.

```
In [29]: # Broadcasting example x1 = torch.rand(5, 5) x2 =
torch.rand(5)      print("Tensor      x1:\n",      x1)
print("Tensor x2:\n", x2) print("x1 - x2:\n", x1 -
x2) print("x1 raised to the power of x2:\n", x1 **
x2)

# Sum of tensor elements along dimension 0
sum_x = torch.sum(x, dim=0)
print("Sum along dimension 0:", sum_x)

# Maximum and minimum values
value, indices = torch.max(x, dim=0)
print("Max value and index:", value, indices)

value, indices = torch.min(x, dim=0)
print("Min value and index:", value, indices)

# Other operations print("Absolute values:", torch.abs(x))
print("Argmax:", torch.argmax(x, dim=0)) print("Argmin:",
torch.argmin(x, dim=0)) print("Mean (converted to float):",
torch.mean(x.float(), dim=0)) print("Element-wise equality (x ==
y):", torch.eq(x, y))

# Sorting
sorted_y, indices = torch.sort(y, dim=0, descending=False)
print("Sorted y and indices:", sorted_y, indices)

# Clamping values
print("Clamped x:", torch.clamp(x, min=0))

# Boolean operations x_bool = torch.tensor([1, 0, 1, 1,
1], dtype=torch.bool) print("Any      True:",
torch.any(x_bool)) print("All True:", torch.all(x_bool))
```



```

Tensor x1:
tensor([[0.5610, 0.6928, 0.8066, 0.2603, 0.1528],
        [0.9867, 0.1102, 0.4665, 0.1929, 0.6669],
        [0.4671, 0.0768, 0.6585, 0.5024, 0.8904],
        [0.6634, 0.6646, 0.0860, 0.1698, 0.0833],
        [0.5990, 0.0964, 0.4688, 0.3539, 0.0450]])

Tensor x2:
tensor([0.8125, 0.2102, 0.6983, 0.5526, 0.1509])

x1 - x2:
tensor([[ -0.2516, 0.4826, 0.1083, -0.2923, 0.0019],
        [0.1607, 0.1102, 0.0390, 0.0000, -0.3547],
        [-0.2341, -0.3547, -0.0503, -0.0503, -0.6123],
        [0.0384, -0.0634, -0.0327, 0.0503, -0.2294],
        [-0.1988, -0.1058, 0.0000, 0.0000, 0.0000]])

x1 raised to the power of x2:
tensor([[0.6252, 0.9258, 0.8606, 0.4753, 0.7532],
        [0.9892, 0.6290, 0.5872, 0.4028, 0.9407],
        [0.5388, 0.5831, 0.7470, 0.6836, 0.9826],
        [0.7164, 0.9177, 0.1803, 0.3754, 0.6873],
        [0.6594, 0.6116, 0.5892, 0.5632, 0.6264]])

Sum along dimension 0: tensor(6) Max value and index:
tensor(3) tensor(2) Min value and index: tensor(1) tensor(0)
Absolute values: tensor([1, 2, 3]) Argmax: tensor(2) Argmin:
tensor(0) Mean (converted to float): tensor(2.) Element-wise
equality (x == y): tensor([False, False, False]) Sorted y and
indices: tensor([7, 8, 9]) tensor([2, 1, 0]) Clamped x:
tensor([1, 2, 3]) Any True: tensor(True) All True:
tensor(False)

```

Tensor Indexing

Access and modify tensor elements using indexing, slicing, and advanced indexing.

- **Accessing Rows & Columns:** Use `x[row, :]` for a row and `x[:, col]` for a column.
- **Slicing:** `x[row, start:end]` extracts a portion of a row.
- **Modifying Elements:** Directly assign values using `x[row, col] = value`.
- **Fancy Indexing:** Use a list of indices to select multiple elements at once.
- **Conditional Indexing:** Extract elements using conditions like `(x < 2) | (x > 8)`.
- **Finding Even Numbers:** Use `x remainder(2) == 0` to filter even values.
- **Conditional Selection with `torch.where()`:** Chooses values based on a condition.

```

In [30]: # Create a random tensor with shape (batch_size, features)
batch_size = 10
features = 25

```

```

x = torch.rand((batch_size, features))

# Access the first row
print("First row of tensor:", x[0, :])

# Access the second column
print("Second column of tensor:", x[:, 1])

# Access the first 10 elements of the third row
print("First 10 elements of third row:", x[2, 0:10])

# Modify a specific element (set first element to 100)
x[0, 0] = 100

# Fancy indexing example
x1 = torch.arange(10)
indices = [2, 5, 8]
print("Fancy indexing result:", x1[indices])

# Advanced indexing: select elements based on a condition
x2 = torch.arange(10)
print("Elements where x2 < 2 or x2 > 8:", x2[(x2 < 2) | (x2 > 8)])
print("Even numbers in x2:", x2[x2.remainder(2) == 0])

# Using torch.where to select values based on a condition
print("Using torch.where:", torch.where(x2 > 5, x2, x2 * 2))

```

```

First row of tensor: tensor([0.1933, 0.0269, 0.3945, 0.6182, 0.3705, 0.7060,
0.4922, 0.0280, 0.3398,
      0.9600, 0.2417, 0.8861, 0.1833, 0.0985, 0.2710, 0.6410, 0.3799, 0.59
81,
      0.0205, 0.9136, 0.9481, 0.6899, 0.9450, 0.6970, 0.1787])
Second column of tensor: tensor([2.6887e-02, 8.1319e-02, 9.8993e-01, 4.5033e
-01, 5.7220e-04, 9.5527e-01,
      1.1555e-01, 9.4050e-02, 5.3863e-02, 5.0582e-01])
First 10 elements of third row: tensor([0.2450, 0.9899, 0.0546, 0.4938, 0.74
71, 0.5465, 0.0106, 0.3488, 0.2002,
      0.4488])
Fancy indexing result: tensor([2, 5, 8])
Elements where x2 < 2 or x2 > 8: tensor([0, 1, 9])
Even numbers in x2: tensor([0, 2, 4, 6, 8])
Using torch.where: tensor([ 0, 2,  4, 6,  8, 10,  6,  7,  8,  9])

```

Tensor Reshaping

Learn how to reshape tensors, concatenate them, and change the order of dimensions.

- **Reshape with `view()` & `reshape()`:** Change tensor shape without altering data.
- **Transpose & Flatten:** `.t()` transposes, `.contiguous().view(-1)` flattens.
- **Concatenation:** `torch.cat([x1, x2], dim=0/1)` merges tensors along rows/columns.

- **Flattening:** `.view(-1)` converts a tensor into a 1D array.
- **Batch Reshaping:** `.view(batch, -1)` keeps batch size while reshaping.
- **Permute Dimensions:** `.permute(0, 2, 1)` reorders dimensions efficiently.
- **Unsqueeze for New Dimensions:** `.unsqueeze(dim)` adds singleton dimensions.

```
In [31]: # Reshape a tensor using view and reshape
x = torch.arange(9)
x_3x3 = x.view(3, 3)
print("Reshaped to 3x3 using view:\n", x_3x3)
x_3x3 = x.reshape(3, 3)
print("Reshaped to 3x3 using reshape:\n", x_3x3)

# Transpose and flatten the tensor
y = x_3x3.t()
print("Flattened transposed tensor:", y.contiguous().view(9))

# Concatenation example x1 = torch.rand(2, 5) x2 = torch.rand(2, 5)
print("Concatenated along dimension 0 (rows):", torch.cat([x1, x2], dim=0).shape)
print("Concatenated along dimension 1 (columns):", torch.cat([x1, x2], dim=1).shape)

# Flatten the tensor using view(-1)
z = x1.view(-1)
print("Flattened tensor shape:", z.shape)

# Reshape with batch dimension
batch = 64
x = torch.rand(batch, 2, 5)
print("Reshaped to (batch, -1):", x.view(batch, -1).shape)

# Permute dimensions
z = x.permute(0, 2, 1)
print("Permuted tensor shape:", z.shape)

# Unsqueeze examples (adding new dimensions) x = torch.arange(10)
print("Original x:", x) print("x unsqueezed at dim 0:",
x.unsqueeze(0).shape, x.unsqueeze(0)) print("x unsqueezed at dim 1:",
x.unsqueeze(1).shape, x.unsqueeze(1))
```



Reshaped to 3x3 using view:

```
tensor([[0, 1, 2],  
        [3, 4, 5],  
        [6, 7, 8]])
```

Reshaped to 3x3 using reshape:

```
tensor([[0, 1, 2],  
        [3, 4, 5],  
        [6, 7, 8]])
```

Flattened transposed tensor: tensor([0, 3, 6, 1, 4, 7, 2, 5, 8]) Concatenated along dimension 0 (rows): torch.Size([4, 5]) Concatenated along dimension 1 (columns): torch.Size([2, 10]) Flattened tensor shape: torch.Size([10]) Reshaped to (batch, -1): torch.Size([64, 10]) Permuted tensor shape: torch.Size([64, 5, 2]) Original x: tensor([0, 1, 2, 3, 4, 5, 6, 7, 8, 9]) x unsqueezed at dim 0: torch.Size([1, 10]) tensor([[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]]) x unsqueezed at dim 1: torch.Size([10, 1]) tensor([[0],

```
[1],  
[2],  
[3],  
[4],  
[5],  
[6],  
[7],  
[8],  
[9]])
```



SUBSCRIBE



AwesomeNeuron

awesomeneuron.substack.com

Found this helpful?



Repost it
to help others
in your network.



Rohit Kumar Tiwari