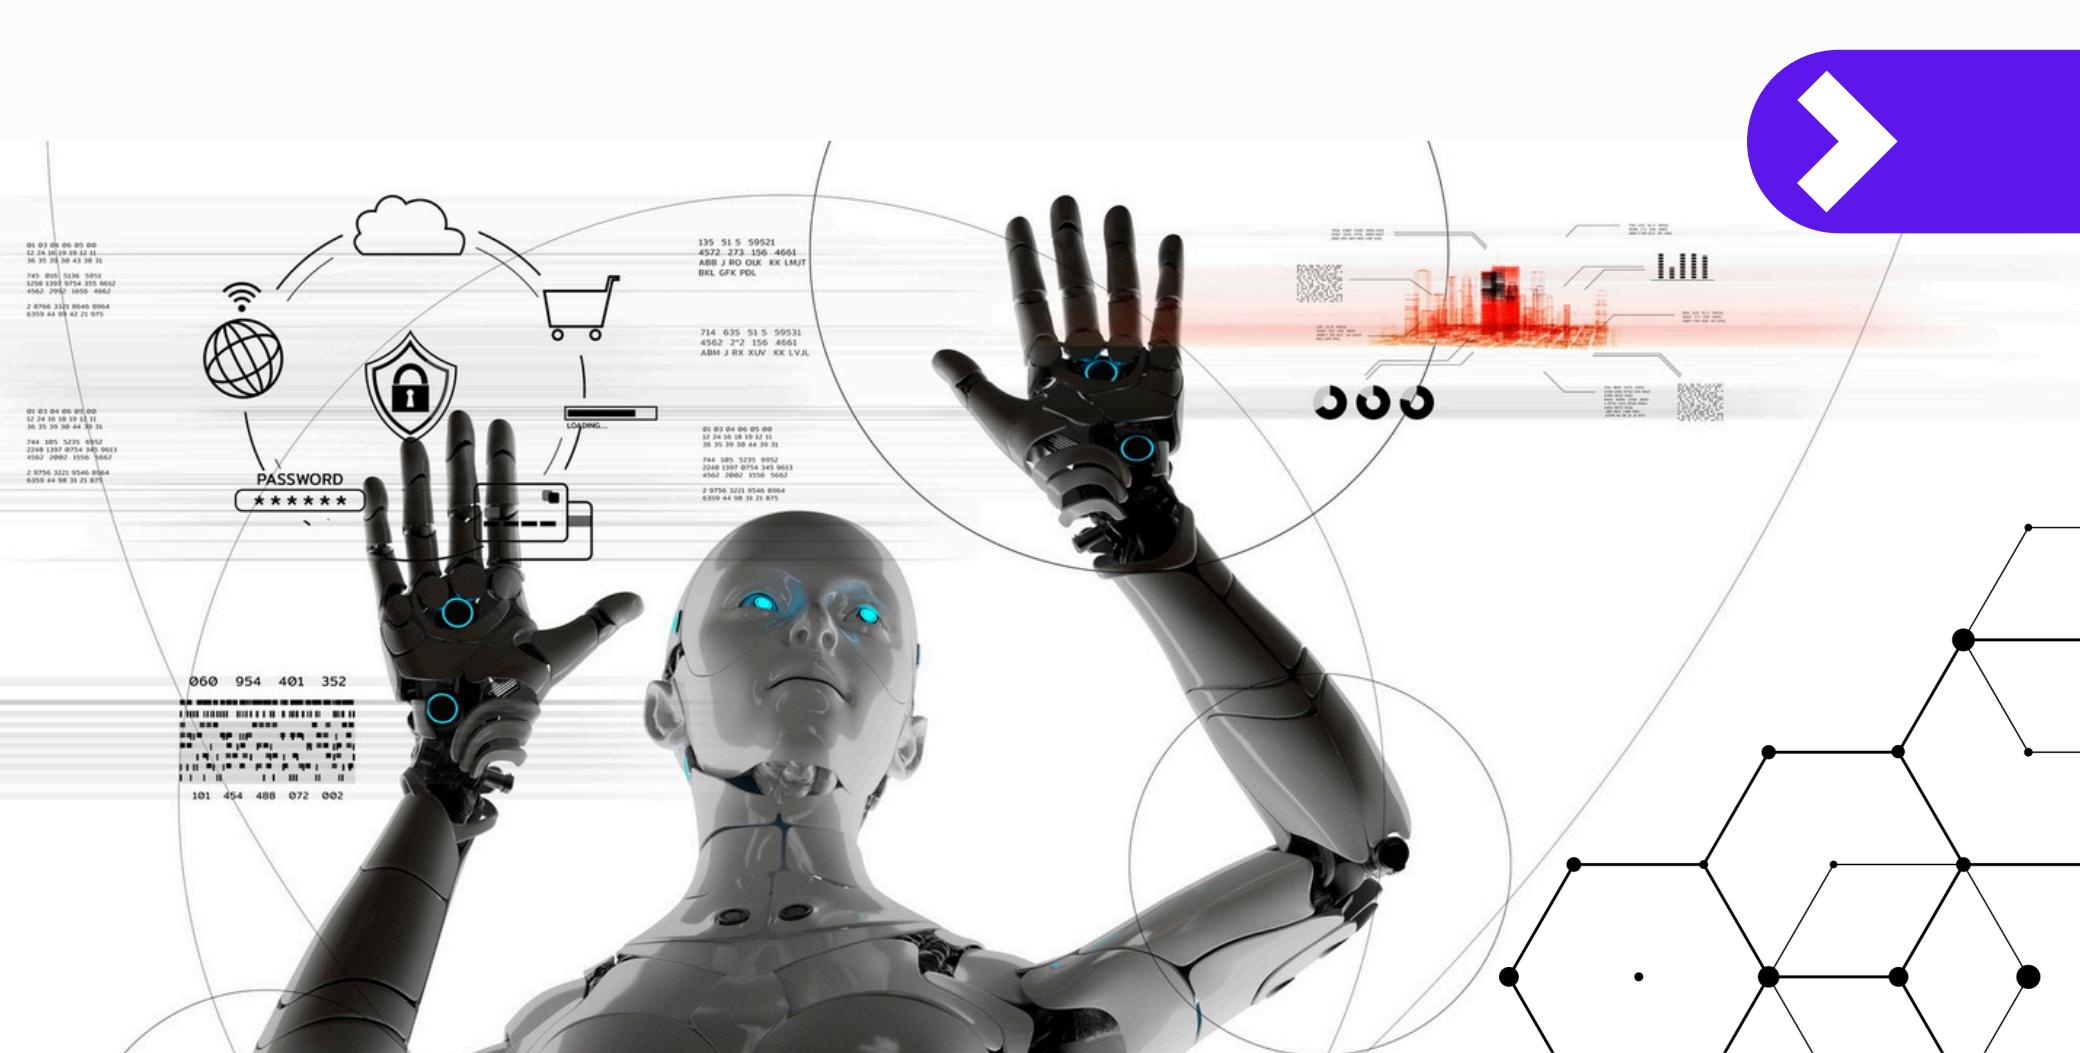




Run LLMs Locally with Ollama



Run powerful LLMs locally and integrate with frameworks like **LangChain & Llamaindex** easily!



What is Ollama?

Your Local AI Powerhouse Ollama is a lightweight, extensible framework that lets you run large language models locally on your machine. No cloud dependencies, no API costs, complete privacy!

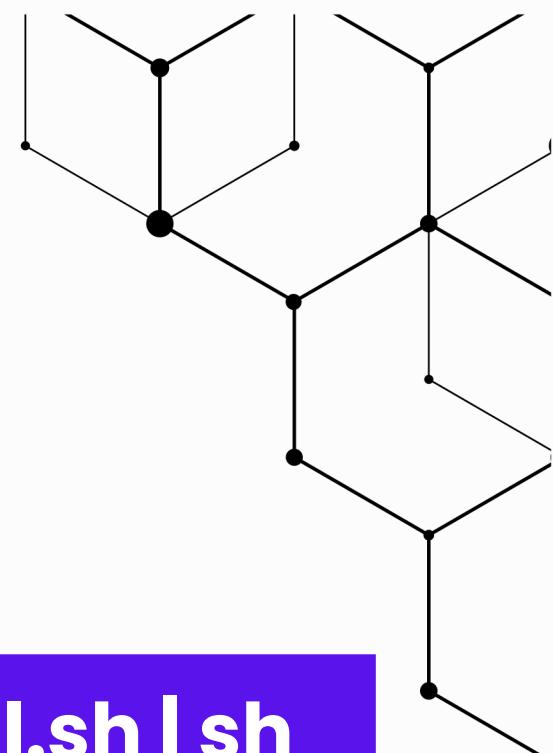
Key Features

- Run models like Mistral, Llama, CodeLlama locally
- Easy installation and setup
- Multiple framework integrations
- Complete data privacy
- Offline capabilities

Perfect for

- Building RAG applications
- Creating AI agents
- Prototyping without API costs
- Privacy-sensitive projects

Installation Process



Quick Installation Steps

Linux:

```
curl -fsSL https://ollama.com/install.sh | sh
```

Download installer from ollama.com

macOS

Download Ollama



Download for macOS

Requires macOS 12 Monterey or later

Windows

Download Ollama



Download for Windows

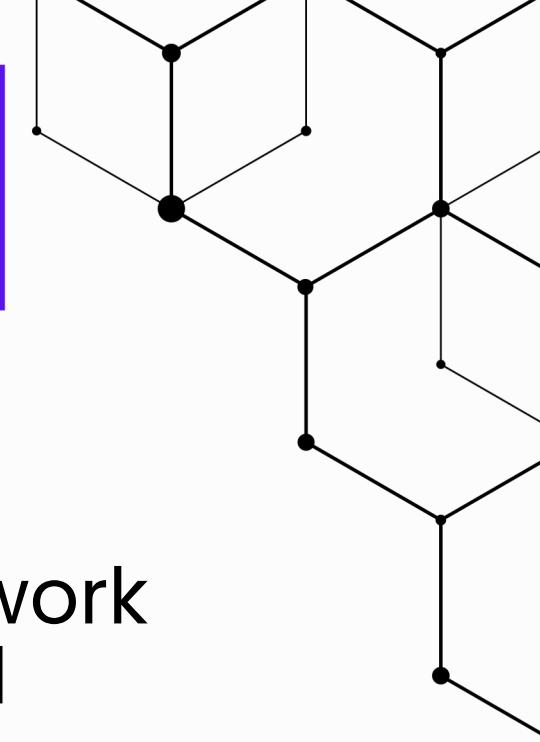
Requires Windows 10 or later

Commands



```
ollama --version      # Check installed Ollama version
ollama serve          # Start Ollama service
ollama pull mistral   # Download Mistral model
ollama pull llama2    # Download LLaMA2 model
ollama pull codellama  # Download CodeLLaMA model
ollama run mistral    # Test installation by running Mistral model
```

Framework Integration



Multiple Integration Options

Supported Frameworks:

- LangChain – Full-featured LLM framework
- LlamaIndex – Document indexing and retrieval
- Haystack – Enterprise-grade NLP pipelines
- Direct API – Pure HTTP requests

Installation Commands:

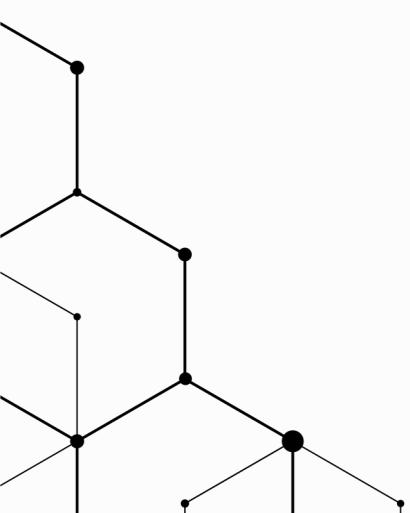


```
# LangChain
pip install langchain-community

# LlamaIndex
pip install llama-index-llms-ollama

# Haystack
pip install ollama-haystack

# Direct API (requests)
pip install requests
```



LangChain Integration



LangChain + Ollama

Code Example:

```
from langchain_community.llms import Ollama

# Initialize Ollama LLM
llm = Ollama(model="mistral")

# Run inference
response = llm("Explain the difference between transformers and RNNs.")

print(response)
```

Benefits

- Seamless integration with LangChain ecosystem
- Access to chains, agents, and tools
- Perfect for RAG applications
- Memory management built-in

Use Cases

- Chatbots with memory
- Document Q&A systems
- Multi-step reasoning tasks

Llamaindex Integration



Llamaindex + Ollama

Code Example:



```
from llama_index.llms.ollama import Ollama

# Initialize with custom settings
llm = Ollama(model="mistral", request_timeout=120.0, context_window=8000)

# Simple inference
response = llm.complete("Summarise India's independence movement in one paragraph.")

print(response)
```

Benefits

- Optimized for document indexing
- Built-in RAG capabilities
- Efficient context management
- Easy data ingestion

Use Cases

- Knowledge base systems
- Document search and retrieval
- Research assistants

Haystack Integration

Haystack + Ollama

Code Example:

```
from haystack_integrations.components.generators.ollama import OllamaChatGenerator
from haystack.dataclasses import ChatMessage

generator = OllamaChatGenerator( model="mistral", url="http://localhost:11434",
                                 generation_kwargs={ "num_predict": 100, "temperature": 0.9, })

messages = [ ChatMessage.from_system("You are a helpful assistant"),
             ChatMessage.from_user("What's Natural Language Processing?")]

print(generator.run(messages=messages))
```

Benefits

- Enterprise-grade pipelines
- Advanced preprocessing
- Scalable architecture
- Production-ready

Direct API Integration

🔌 Direct API Access

Code Example:

```
import requests

url = "http://localhost:11434/api/generate"

prompt = "What is the capital of France?"

response = requests.post(url, json={"model": "mistral", "prompt": prompt, "stream": False})

print(response.json()['response'])
```

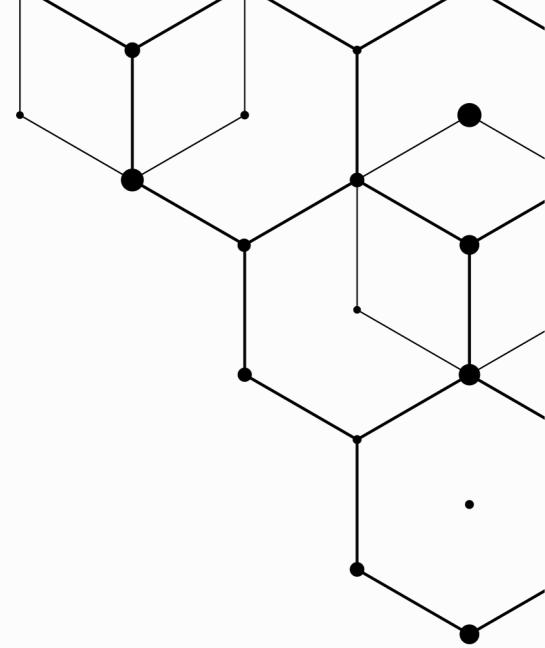
Benefits

- Maximum control and flexibility
- Lightweight integration
- Custom request handling
- No additional dependencies

Use Cases

- Custom applications
- Microservices
- Simple integrations

Benefits & Advantages



✓ Why Choose Ollama?

Privacy & Security:

- Data never leaves your machine
- No cloud dependencies
- Complete control over your data
- GDPR/compliance friendly

Cost Efficiency:

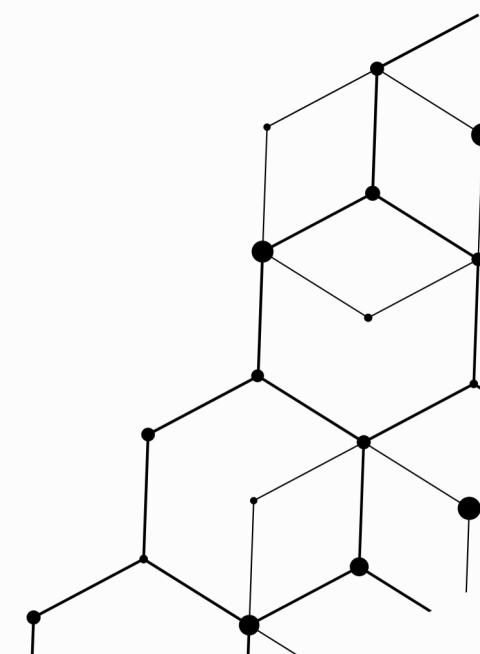
- No API fees or usage limits
- One-time setup cost
- Unlimited inference
- Perfect for experimentation

Performance:

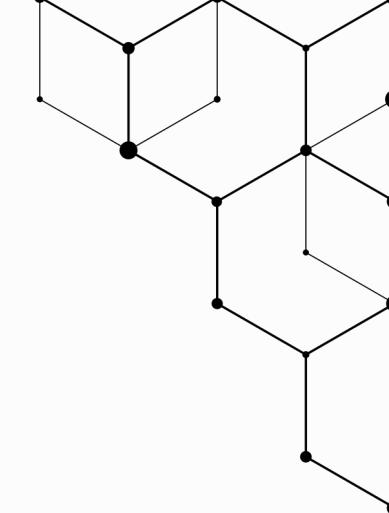
- Low latency responses
- Offline capabilities
- Consistent performance
- Hardware optimization

Flexibility:

- Multiple model choices
- Custom model fine-tuning
- Various framework integrations
- Open-source ecosystem



Popular Models



Choose the Right Model

deepseek-r1

DeepSeek-R1 is a family of open reasoning models with performance approaching that of leading models, such as Qwen and Gemini 2.5 Pro.

tools thinking 1.5b 7b 8b 14b 32b 70b 67b

↓ 52.1M Pulls ⚡ 35 Tags ⏰ Updated 3 days ago

gemma3n

Gemma 3n models are designed for efficient execution on everyday devices such as laptops, tablets or phones.

e2b e4b

↓ 74.4K Pulls ⚡ 9 Tags ⏰ Updated 1 week ago

gemma3

The current, most capable model that runs on a single GPU.

vision 1b 4b 12b 27b

↓ 7.7M Pulls ⚡ 21 Tags ⏰ Updated 2 months ago

qwen3

Qwen3 is the latest generation of large language models in Qwen series, offering a comprehensive suite of dense and mixture-of-experts (MoE) models.

tools thinking 0.6b 1.7b 4b 8b 14b 30b 32b 235b

↓ 3.3M Pulls ⚡ 35 Tags ⏰ Updated 1 month ago

qwen2.5vl

Flagship vision-language model of Qwen and also a significant leap from the previous Qwen2-VL.

vision 3b 7b 32b 72b

↓ 332.3K Pulls ⚡ 17 Tags ⏰ Updated 1 month ago

llama3.1

LLama 3.1 is a new state-of-the-art model from Meta available in 8B, 70B and 405B parameter sizes.

tools 8b 70b 405b

↓ 97M Pulls ⚡ 93 Tags ⏰ Updated 7 months ago

llama3.2

Meta's Llama 3.2 goes small with 1B and 3B models.

tools 1b 3b

↓ 23.4M Pulls ⚡ 63 Tags ⏰ Updated 9 months ago

mistral

The 7B model released by Mistral AI, updated to version 0.3.

tools 7b

↓ 16.2M Pulls ⚡ 84 Tags ⏰ Updated 2 weeks ago

qwen2.5

Qwen2.5 models are pretrained on Alibaba's latest large-scale dataset, encompassing up to 18 trillion tokens. The model supports up to 128K tokens and has multilingual support.

tools 0.5b 1.5b 3b 7b 14b 32b 72b

↓ 10.8M Pulls ⚡ 133 Tags ⏰ Updated 9 months ago

llama3

Meta Llama 3: The most capable openly available LLM to date

8b 70b

↓ 9M Pulls ⚡ 68 Tags ⏰ Updated 1 year ago

llava

LLava is a novel end-to-end trained large multimodal model that combines a vision encoder and Vicuna for general-purpose visual and language understanding. Updated to version 1.6.

vision 7b 13b 34b

↓ 7.4M Pulls ⚡ 98 Tags ⏰ Updated 1 year ago

gemma2

Google Gemma 2 is a high-performing and efficient model available in three sizes: 2B, 9B, and 27B.

2b 9b 27b

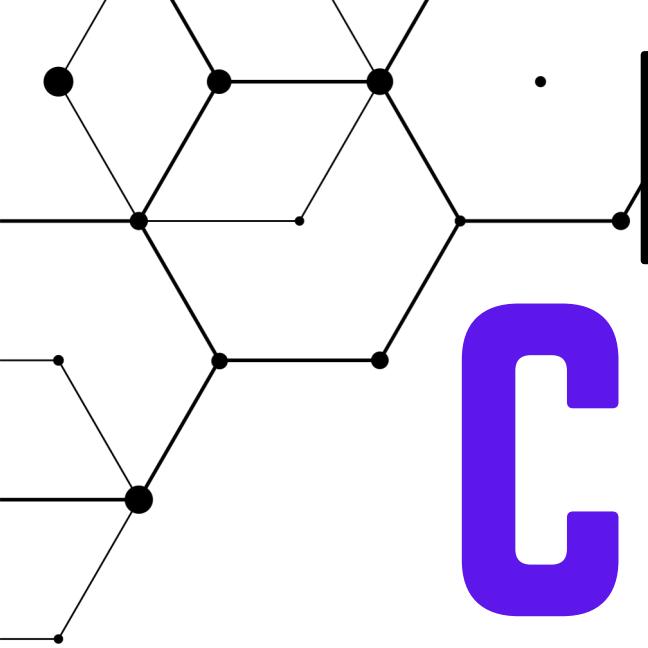
↓ 5.7M Pulls ⚡ 94 Tags ⏰ Updated 11 months ago

qwen2.5-coder

The latest series of Code-Specific Qwen models, with significant improvements in code generation, code reasoning, and code fixing.

tools 0.5b 1.5b 3b 7b 14b 32b

↓ 5.7M Pulls ⚡ 199 Tags ⏰ Updated 1 month ago



LIKE THIS CONTENT ?

FOLLOW FOR MORE!



 **LIKE**

 **REPOST**

 **SAVE**

