




# Popular AI Agent Frameworks



@rakeshgoehel01






# Brief

**Agent frameworks link agent components like nerves in a body**

That's why choosing the right one can make or break your entire agent ecosystem.

Today, we will take a look at each of these framework to identify which fits your use cases better. →



# LangGraph

**LangChain is an established framework in the agentic AI ecosystem, and LangGraph extends its capabilities to support complex, stateful agent workflows.**

## ■ KEY FEATURES OF LANGGRAPH

- **Component ecosystem:** They provide an extensive library of pre-built components that enable the rapid development of specialized agents.
- **Foundation model selection:** The frameworks support diverse foundation models like Anthropic Claude and Amazon Nova models, offering different reasoning capabilities.
- **Graph-based workflows:** LangGraph allows for the definition of complex agent behaviors as state machines, supporting sophisticated decision logic.
- **Memory abstractions:** They offer multiple options for short-term and long-term memory management, which is essential for agents that need to maintain context over time.

# Google ADK

**Google's Agent Development Kit (ADK) is an open-source framework designed to simplify the full-stack end-to-end development of agents and multi-agent systems.**

## KEY FEATURES OF GOOGLEADK

- **Model-Agnostic Design:** While optimized for Gemini models, ADK can work with any LLM provider.
- **Flexible Orchestration:** Define workflows using structured workflow agents (Sequential, Parallel, Loop) for predictable pipelines, or leverage LLM-driven dynamic routing for adaptive behavior.
- **Multi-Agent Architecture:** Build modular and scalable applications by composing multiple specialized agents.
- **Rich Tool Ecosystem:** Equip agents with diverse capabilities using pre-built tools (Search, Code Execution), create custom functions, integrate third-party libraries (LangChain, CrewAI), or use other agents as tools.



**CrewAI is an open-source framework focused on autonomous multi-agent orchestration. It provides a structured approach for creating teams of specialized autonomous agents.**

## KEY FEATURES OF CREWAI

- **Role-based agent design:** Autonomous agents are defined with specific roles, goals, and backstories to enable specialized expertise.
- **Task delegation:** The framework includes built-in mechanisms for autonomously assigning tasks to the most appropriate agents based on their capabilities.
- **Agent collaboration:** It provides a framework for autonomous inter-agent communication and knowledge sharing without human mediation.
- **Process management:** CrewAI offers structured workflows for both sequential and parallel autonomous task execution



# Agents SDK

**The OpenAI Agents SDK is a lightweight, Python-first package for building agentic AI applications with minimal abstractions. It is a production-ready evolution of OpenAI's earlier experiments OpenAI Swarm.**

## ■ KEY FEATURES OF AGENTS SDK


- **Python-First Design:** It encourages using native Python features for orchestration and chaining agents, reducing the need to learn new, complex abstractions.
- **Agent Loop:** Provides a built-in agent loop that handles the logic of calling tools, sending the results back to the LLM, and continuing until a task is complete.
- **Handoffs:** A key feature that allows one agent to delegate specific tasks to other, more specialized agents, enabling multi-agent collaboration.
- **Guardrails:** Enables the validation of inputs to agents, allowing checks to run in parallel and fail early if they don't meet criteria.

# AG Autogen

**Initially released by Microsoft, AutoGen is an open-source framework that focuses on enabling conversational and collaborative autonomous AI agents. It provides a flexible, event-driven architecture for building multi-agent systems.**

## KEY FEATURES OF AUTOGEN

- **Conversational agents:** The framework is built around natural language conversations between autonomous agents, facilitating sophisticated reasoning.
- **Asynchronous architecture:** Its event-driven design supports non-blocking agent interactions and complex parallel workflows.
- **Human-in-the-loop:** AutoGen provides strong support for optional human participation in otherwise autonomous agent workflows.
- **Code generation and execution:** It offers specialized capabilities for autonomous agents that can write and run code.




# Semantic Kernel

**Semantic Kernel is an open-source SDK from Microsoft designed to integrate large language models with conventional programming languages like C# and Python**

## KEY FEATURES OF SEMANTIC KERNEL

- **Plugin Architecture:** Its core feature is the ability to expose native code (e.g., C# functions) to an LLM as reusable plugins that the model can call to perform actions.
- **Planner:** Semantic Kernel includes a "planner" that can automatically generate a multi-step plan to achieve a user's goal by combining different plugins.
- **Memory:** It provides built-in capabilities for storing and retrieving conversational history and other contextual information, giving agents long-term memory.
- **Connectors:** The framework offers native integration with OpenAI, Azure OpenAI, and Hugging Face models, making it flexible in its choice of LLMs.



# Microsoft Agent Framework

**Microsoft's unified, production-ready SDK and runtime that merges Semantic Kernel's enterprise-grade plugins, memory, and orchestration with AutoGen.**

## KEY FEATURES OF AGENT FRAMEWORK

- **Graph-based workflows** with typed edges, branching, parallelism, and checkpointing.
- **First-class multi-agent patterns:** sequential, concurrent, handoffs, group chat, Magentic.
- **Open standards:** MCP, OpenAPI, and Agent-to-Agent interoperability.
- **Enterprise ops:** observability, approvals, durability, and recovery for long runs.
- **Strong typing and validation** across agents and workflows to reduce runtime errors.



**Strands Agents is an open-source SDK, initially released by AWS, designed with a model-first approach for building autonomous AI agents.**

## KEY FEATURES OF AWS STRANDS

- **Model-first design:** The framework is built around the concept that the foundation model is the core of agent intelligence, which allows for sophisticated autonomous reasoning.
- **MCP integration:** It has native support for the Model Context Protocol (MCP), enabling standardized context provision to LLMs for consistent operation.
- **AWS service integration:** Strands Agents provides seamless connections to Amazon Bedrock, AWS Lambda, AWS Step Functions, and other services for comprehensive autonomous workflows.
- **Foundation model selection:** It supports a variety of foundation models, including Anthropic Claude and Amazon Nova models on Amazon Bedrock, to optimize for different reasoning capabilities.



# Pydantic Agents

**Pydantic AI is a Python agent framework built on the highly popular Pydantic data validation library. It is designed to build production-grade applications by focusing on robust data validation, and schema enforcement**

## KEY FEATURES OF PYDANTIC AGENTS

- **Data Validation and Parsing:** Its core strength lies in rigorous data validation and type safety, ensuring that data structures passed to and from LLMs are consistent and error-free.
- **Structured Outputs:** Pydantic AI excels at coercing unstructured LLM outputs into structured Pydantic models, which is crucial.
- **Schema Enforcement:** It automates the enforcement of data schemas, which is critical when working with structured APIs and data pipelines in an agentic context.
- **Seamless Integration:** It integrates smoothly with other agent frameworks and is often used alongside them to handle the data layer.



# LLama Index

**LlamalIndex is a data framework for LLM applications that specializes in connecting custom data sources to large language models. It provides tools for data ingestion, indexing, querying and a lot others.**

## ■ KEY FEATURES OF LLAMAINDEX

- **Data Ingestion and Indexing:** LlamalIndex excels at connecting to and ingesting data from various sources (PDFs, APIs, databases) and indexing it for efficient retrieval.
- **Query Engine:** It offers a powerful query engine that allows natural language questions over your data, abstracting away the complexity of vector database queries.
- **RAG Capabilities:** The framework is purpose-built for creating advanced RAG pipelines, enabling LLMs to answer questions based on private or custom data.
- **Agentic Features:** LlamalIndex can be used to build knowledge agents that can perform automated reasoning and decision-based on a given knowledge base and task list.



# Haystack

**Haystack is an open-source framework for building end-to-end AI applications that utilize large language models. It is highly modular and flexible, allowing developers to create custom pipelines for a wide range of GenAI Tasks.**

## KEY FEATURES OF HAYSTACK

- **Pipeline-Based Architecture:** Haystack allows developers to build custom data processing and query pipelines by combining different nodes.
- **Advanced RAG:** It supports building high-performing Retrieval-Augmented Generation pipelines with various strategies, including hybrid retrieval and self-correction loops.
- **Agentic Pipelines:** The framework supports complex workflows using the function-calling capabilities of LLMs, enabling branching and looping for multi-step agentic tasks.
- **Multimodal AI:** Haystack can be used to develop applications that process not just text but also other modalities.



**The Bee Agent Framework is an open-source toolkit developed by IBM Research for building, deploying, and managing scalable agent-based workflows. It is designed to simplify the development of distributed agent systems.**

## KEY FEATURES OF BEE AGENT FRAMEWORK

- **Scalable Workflows:** The framework is designed to construct and serve powerful agentic workflows at scale, handling complex and distributed tasks.
- **Multi-Agent Systems:** Recent updates have added extensions for multi-agent collaboration, allowing multiple "bee" agents to work together to solve a problem
- **Developer Experience:** A key goal of the framework is to "make simple things simple and complex things possible," with a focus on ease of use.
- **Observability:** The ecosystem includes tools like Bee UI for visual interaction and Bee Observe for telemetry collection and management



# Smol Agents

Hugging Face's **smolagents** is a Python library designed to simplify the creation of AI agents, making agentic AI more accessible to a broader range of developers. It aims to be a plug-and-play framework that requires minimal setup.

## KEY FEATURES OF SMOL AGENTS

- **Simplicity:** The library is designed for ease of use, allowing a developer to create and run a functional agent in just a few lines of code.
- **Model Agnostic:** While part of the Hugging Face ecosystem, it can work with various models, including those from the Hugging Face Hub, OpenAI, and Anthropic, as well as local models.
- **Modality-Agnostic:** Beyond text, smolagents can handle vision, video, and audio inputs, which broadens the range of possible applications.
- **Custom Tools:** The framework makes it straightforward to build custom tools for an agent to use, alongside built-in tools like web search.



# Agno (Prev. Phidata)

**Previously Known as Phidata. Agno is an open-source, high-performance Python framework for multi-agent workflows with a built-in FastAPI runtime and control plane.**

## KEY FEATURES OF SMOL AGENTS

- **Unified, Pythonic framework** for single agents, teams, and step-based workflows (sequential, parallel, branching, loops).
- **Ready-to-use FastAPI “AgentOS”** app for serving agents with a built-in control plane for testing, monitoring, and management.
- **High performance and low overhead** focus with async runtime, minimal memory footprint, and horizontal scalability.
- **Declarative agent composition:** configure model, memory, tools, and knowledge sources with simple Python.
- **Transparent reasoning and observability:** inspect traces, tool calls, and logs for reliability and auditability.

# Framework Use Cases

## Langgraph

- Suited for complex, multi-step reasoning with rich component orchestration for autonomous decision-making.
- Great for Python teams needing large plugin ecosystems and robust memory/state management.

## Google ADK

- Optimal for developing production-grade, multi-agent applications needing flexible orchestration and modality support.
- Recommended for teams requiring seamless interoperability, advanced streaming, or tight Google Cloud integration.

## CrewAI

- Effective for problems benefitting from specialized, role-based agent teams autonomously collaborating.
- Useful when explicit multi-agent coordination and division of labor enhance solution quality.

## OpenAI's Agents SDK

- Ideal for quickly building agentic applications tightly integrated with OpenAI's ecosystem using Python.
- Great for action-oriented AI apps that must call APIs, delegate tasks, or use a lightweight agent/control loop.

## Autogen

- Fits applications desiring natural language, conversational agent flows and collaborative multi-agent reasoning.
- Good for scenarios combining autonomous operation with optional human oversight or code generation/execution.

## Semantic Kernel

- Well-suited for integrating LLMs with business logic via code/plugin functions in enterprise applications.
- Good for building AI copilots and automated workflows

## Microsoft Agent Framework

- Compose specialized agents (data retrieval, analysis, policy/compliance) that collaborate via structured handoffs
- Run long-lived, stateful processes that autonomously execute tools and invite approvals

## AWS Strands

- Ideal for teams leveraging AWS infrastructure needing deep, native integration with AWS services.
- Best for enterprise applications demanding flexibility

## Pydantic AI

- Best as a supporting layer for any agentic/ML application where rigorous data validation, and type safety are critical.
- Useful in AI APIs and data pipelines needing structured, predictable inputs/outputs from LLMs.

## Llama Index

- Best for building applications powered by custom/unstructured data, including RAG systems.
- Ideal when you need agents or chatbots to provide natural language access to large, private or structured data sources.

## Haystack

- Excellent for creating flexible, composable LLM pipelines
- Favored in use cases demanding multimodal support, large-scale document search, or complex agentic workflows.

## IBM Bee

- Suited for building robust, distributed multi-agent systems
- Ideal for enterprise scenarios orchestrating multiple specialized agents

## HF Smol Agents

- Best for developers needing fast, simple prototyping
- Useful in education, experimentation, or rapid tool-driven agents

## Agno (Prev. Phidata)

- Private, in-cloud AgentOS runtime with built-in control plane for testing
- High-performance, step-based multi-agent workflows that are stateless, async, and horizontally scalable

# Conclusion

---

**When selecting an agentic AI framework for autonomous agent development,**

**Consider how each option aligns with your specific requirements, including technical capabilities, organizational fit, team expertise, existing infrastructure, and long-term maintenance requirements.**

*If you want to learn more,*

*Each framework's repo/documentation is attached in the comment section of this document's post.*

Hi, I am  
Rakesh Gohel



**“We help businesses 10X their growth with  
Cloud and AI Agents”**

**FOLLOW TO LEARN MORE ABOUT AI AGENTS**

[linkedin.com/in/rakeshgohel01](https://linkedin.com/in/rakeshgohel01)